

# AUDITING SCALE-LAW EXCEPTIONS IN ROBOTICS

**Anonymous authors**

Paper under double-blind review

## ABSTRACT

Scaling laws predict broad improvements in model loss as data, parameters, and compute increase. Robot policies, however, are deployed at physical decision boundaries: rare contacts, hidden preconditions, morphology shifts, spurious visual shortcuts, and recovery-scarce failures. We rebuild Paper 107 around a falsifiable claim: a robot policy can follow an average scaling trend while still crossing the wrong embodied failure boundary. The local benchmark covers 5 robot task families, 7 scale-exception regimes, 5 stress splits, 9 methods, 7 seeds, and 84 episodes per group. On combined exception stress, the proposed embodied scale-exception audit obtains  $0.550 \pm 0.005$  success versus  $0.407 \pm 0.006$  for the strongest non-oracle baseline, improves rare-failure recall from 0.375 to 0.512, reduces contact-boundary error from 0.283 to 0.225, reduces unsafe commitment from 0.086 to 0.073, and reduces overtrust from 0.106 to 0.069. Pairwise seed tests favor the proposed method over the strongest baseline in 7/7 seeds, and ablations show that rare-tail memory, contact-boundary factors, embodiment-shift detection, shortcut auditing, and recoverability calibration each matter. A 2026-06-15 continuation audit reran the full benchmark and preserved the terminal decision: **STRONG REVISE**. The local evidence supports the mechanism, but the paper is not ICLR-main-ready until validated on real robots or independent high-fidelity benchmarks.

## 1 MOTIVATION

Scaling laws are one of the clearest empirical regularities in modern machine learning (Kaplan et al., 2020; Hoffmann et al., 2022). They are also increasingly relevant to robotics. Generalist agents, robot transformers, vision-language-action models, cross-embodiment datasets, and open robot-policy models show that broader training can produce more capable robot policies (Reed et al., 2022; Brohan et al., 2022; 2023; Open X-Embodiment Collaboration, 2023; Octo Model Team, 2024).

But robot deployment is not only an average-loss problem. A policy that improves average imitation accuracy can still overtrust visual shortcuts, miss rare contact tails, fail under morphology shift, or commit to an unrecoverable action. In these cases, the scaling curve looks healthy while the embodied failure boundary moves in the wrong direction.

## 2 METHOD SKETCH

The proposed method, `proposed_embodied_scale_exception_audit`, evaluates scaling curves through embodied exception metrics rather than average success alone. It tracks:

- rare physical failure recall;
- contact-boundary error;
- embodiment-shift sensitivity;
- spurious shortcut reliance;
- unsafe commitment and overtrust;
- calibration error under rare tails;
- recoverability after a bad scaled-policy commitment.

Table 1: Combined-stress scale-law exception benchmark.

| Method       | Succ. | RareRec. | ContactErr. | Unsafe | Overtrust | Regret |
|--------------|-------|----------|-------------|--------|-----------|--------|
| Oracle       | 0.710 | 0.626    | 0.180       | 0.031  | 0.022     | 0.000  |
| Proposed     | 0.550 | 0.512    | 0.225       | 0.073  | 0.069     | 0.159  |
| Conformal    | 0.407 | 0.375    | 0.283       | 0.086  | 0.106     | 0.303  |
| EnsembleGate | 0.386 | 0.345    | 0.300       | 0.099  | 0.128     | 0.323  |
| ScaleLaw     | 0.315 | 0.216    | 0.334       | 0.155  | 0.231     | 0.395  |
| DataScale    | 0.295 | 0.227    | 0.331       | 0.151  | 0.198     | 0.415  |
| ModelScale   | 0.290 | 0.189    | 0.345       | 0.157  | 0.233     | 0.420  |
| LargeBC      | 0.275 | 0.165    | 0.359       | 0.167  | 0.241     | 0.435  |
| SmallBC      | 0.201 | 0.183    | 0.351       | 0.159  | 0.153     | 0.509  |

The method does not claim that scaling is bad. Instead, it asks whether a scaling trend is predictive of the boundary that matters for physical deployment. The implementation is an executable diagnostic benchmark rather than a trained robot foundation model, so it supports local mechanism auditing but not final main-track empirical claims.

### 3 BENCHMARK

The rebuilt benchmark is generated by `src/run_experiment.py`. It writes aggregate CSVs, figures, and LaTeX tables directly, preserving full coverage while staying lightweight.

**Tasks.** The tasks are contact insertion, deformable handling, mobile manipulation, bimanual handoff, and tool-use sequence.

**Scale-exception regimes.** The regimes are rare contact tail, occluded precondition, embodiment mismatch, long-horizon compounding, spurious visual shortcut, recovery-scarce failure, and mixed exception stress.

**Splits.** The splits are nominal scaling, contact-tail shift, embodiment shift, shortcut shift, and combined exception stress.

**Methods.** The baselines are small behavior cloning, large behavior cloning, data-only scaling, model-only scaling, generic scaling-law extrapolation, ensemble uncertainty gating, and conformal risk filtering. The benchmark also includes the proposed audit and an oracle exception selector.

**Metrics.** The evidence reports task success, rare-failure recall, contact-boundary error, tail-risk rate, unsafe commitment, overtrust, calibration ECE, recovery success, intervention cost, data-efficiency proxy, regret to oracle, paired seed differences, ablations, stress sweeps, and failure cases.

### 4 RESULTS

The strongest non-oracle baseline is `conformal_risk_filter`. On combined exception stress, the proposed method improves success by  $0.144 \pm 0.006$  in paired seed comparisons, wins 7/7 seeds, improves rare-failure recall by 0.138, reduces contact-boundary error by 0.058, reduces unsafe commitment by 0.013, and reduces overtrust by 0.037. The oracle remains substantially higher at  $0.710 \pm 0.010$ , so the benchmark is not saturated.

At maximum stress level 0.95, the proposed method remains above the strongest non-oracle baseline with  $0.539 \pm 0.007$  success versus  $0.389 \pm 0.009$ , higher rare-failure recall (0.505 versus 0.355), lower contact-boundary error (0.224 versus 0.291), lower tail risk (0.135 versus 0.179), lower unsafe commitment (0.071 versus 0.092), and lower overtrust (0.072 versus 0.111). This is still local evidence rather than external validation.




Figure 1: Combined-exception success with 95% confidence intervals. The proposed embodied scale-exception audit clears all non-oracle baselines but remains below the oracle.




Figure 2: Exception-boundary diagnostics. The proposed method improves rare-failure recall while lowering contact-boundary error under combined exception stress.

## 5 ABLATIONS AND STRESS TESTS

The best removed-component ablation is `minus_embodiment_shift_detector`, which still trails the full audit by 0.051 success. Removing rare-tail memory, contact-boundary factors, shortcut auditing, or recoverability calibration produces distinct diagnostic failures. This supports the claim that the method is not just larger scaling, uncertainty gating, or conformal filtering.

## 6 WHERE THE METHOD STILL FAILS

The failure-case CSV shows that the audit helps least when ordinary uncertainty gating already exposes the exception or when the scaled policy is not overtrusted. Those are intended boundaries.

Table 2: Pairwise combined-exception success differences against the proposed method.

| Baseline     | Diff   | CI    | Wins |
|--------------|--------|-------|------|
| Conformal    | 0.144  | 0.006 | 7    |
| DataScale    | 0.256  | 0.009 | 7    |
| EnsembleGate | 0.164  | 0.005 | 7    |
| ScaleLaw     | 0.235  | 0.009 | 7    |
| LargeBC      | 0.275  | 0.006 | 7    |
| ModelScale   | 0.261  | 0.012 | 7    |
| Oracle       | -0.159 | 0.009 | 0    |
| SmallBC      | 0.349  | 0.006 | 7    |

Table 3: Ablations of the embodied scale-exception audit.

| Ablation       | Succ. | RareRec. | ContactErr. | Unsafe |
|----------------|-------|----------|-------------|--------|
| Full           | 0.554 | 0.515    | 0.220       | 0.072  |
| NoEmbShift     | 0.503 | 0.480    | 0.234       | 0.086  |
| NoShortcut     | 0.485 | 0.473    | 0.237       | 0.091  |
| NoRecoverCalib | 0.483 | 0.483    | 0.240       | 0.103  |
| NoContact      | 0.475 | 0.476    | 0.299       | 0.101  |
| NoRareMem      | 0.464 | 0.307    | 0.256       | 0.083  |
| UncertOnly     | 0.379 | 0.343    | 0.295       | 0.100  |

The method’s value appears when average scaling metrics are misleading about rare contact tails, shortcuts, embodiment mismatch, or scarce recovery budget.

## 7 HOSTILE PRIOR-WORK BOUNDARY

Kaplan-style and Chinchilla-style scaling work already show that model behavior can be predicted as data, parameters, and compute change (Kaplan et al., 2020; Hoffmann et al., 2022). Gato, RT-1, RT-2, RT-X/Open X-Embodiment, and Octo already demonstrate broad robot or embodied policy scaling (Reed et al., 2022; Brohan et al., 2022; 2023; Open X-Embodiment Collaboration, 2023; Octo Model Team, 2024).

The present paper survives only under a narrower interpretation: it is about *embodied scale exceptions*, not scaling laws or robot foundation models in general. If real robot experiments later show that scaled robot policies, conformal filters, or ensemble uncertainty gates match the proposed audit on success, rare-failure recall, contact-boundary error, unsafe commitment, overtrust, and ablation robustness, the claim should be killed.

## 8 LIMITATIONS AND TERMINAL DECISION

The evidence is strong enough to justify continued work but not submission. The limitations are material:

- the benchmark is local and synthetic;
- no real robot or independent high-fidelity simulator has been used;
- baselines are executable diagnostic models, not external trained robot systems;
- no learned policy checkpoint, training curve, or deployment video is provided;
- no external benchmark such as LIBERO, RL Bench, Meta-World, BridgeData, Open X-Embodiment, DROID, or a hardware manipulation suite is reported.

The terminal decision for this rebuild is therefore **STRONG REVISE**. The next honest version must add real robot or independent high-fidelity evidence, implemented learned baselines, and qualitative rollouts. Without that evidence, the paper should not be submitted to ICLR main.




Figure 3: Ablation success under combined exception stress. The full audit remains above all removed-component variants.




Figure 4: Stress sweep over exception intensity. The proposed method remains above non-oracle baselines across the generated stress range.

## REFERENCES

Anthony Brohan, Noah Brown, Justice Carbajal, Yevgen Chebotar, Xi Chen, Krzysztof Choromanski, Tianli Ding, Danny Driess, Avinava Dubey, Chelsea Finn, Pete Florence, Chuyuan Fu, Montse Gonzalez Arenas, Kehang Gopalakrishnan, Kehang Han, Karol Hausman, Alex Herzog, Jasmine Hsu, Brian Ichter, Alex Irpan, Nikhil Joshi, Ryan Julian, Dmitry Kalashnikov, Yuheng Kuang, Isabel Lee, Sergey Levine, Yao Lu, Carolina Parada, Peter Pastor, John Quiambao, Kanishka Rao, Igor Rettinghouse, Diego Reyes, Pierre Sermanet, Nicholas Sievers, Clayton Tan, Alexander Toshev, Vincent Vanhoucke, Fei Xia, Ted Xiao, Peng Xu, Sichun Xu, Mengyuan Yu, and Brianna Zitkovich. Rt-1: Robotics transformer for real-world control at scale, 2022.




Figure 5: Unsafe commitment and regret trade-off. The proposed method improves success while lowering unsafe commitment relative to the strongest non-oracle baseline.

Anthony Brohan, Noah Brown, Justice Carbajal, Yevgen Chebotar, Joseph Dabis, Chelsea Finn, Kehang Gopalakrishnan, Karol Hausman, Alex Herzog, Jasmine Hsu, Brian Ichter, Alex Irpan, Tyler Jackson, Sabeek Jesmonth, Nikhil Joshi, Ryan Julian, Dmitry Kalashnikov, Yuheng Kuang, Isabel Lee, Sergey Levine, Yao Lu, Utsav Malla, Deeksha Manjunath, Igor Mordatch, Ofir Nachum, Carolina Parada, Jodilyn Peralta, Emily Perez, Karl Pertsch, John Quiambao, Kanishka Rao, Michael Ryoo, Grecia Salazar, Pannag Sanketi, Pierre Sermanet, Jaspiar Singh, Clayton Tan, Huong Tran, Vincent Vanhoucke, Quan Vuong, Fei Xia, Ted Xiao, Peng Xu, Sichun Xu, Tianhe Yu, Brianna Zhang, and Brianna Zitkovich. Rt-2: Vision-language-action models transfer web knowledge to robotic control, 2023.

Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, Tom Hennigan, Eric Noland, Katherine Millican, George van den Driessche, Bogdan Damoc, Aurelia Guy, Simon Osindero, Karen Simonyan, Erich Elsen, Jack W. Rae, Oriol Vinyals, and Laurent Sifre. Training compute-optimal large language models, 2022.

Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models, 2020.

Octo Model Team. Octo: An open-source generalist robot policy, 2024.

Open X-Embodiment Collaboration. Open x-embodiment: Robotic learning datasets and rt-x models, 2023.

Scott Reed, Konrad Zolna, Emilio Parisotto, Sergio Gomez Colmenarejo, Alexander Novikov, Gabriel Barth-Maron, Mai Gimenez, Yury Sulsky, Jackie Kay, Jost Tobias Springenberg, Tom Eccles, Jake Bruce, Ali Razavi, Ashley Edwards, Nicolas Heess, Yutian Chen, Raia Hadsell, Oriol Vinyals, Misha Bordbar, and Nando de Freitas. A generalist agent, 2022.