

SCALE-LAW EXCEPTIONS IN ROBOTICS: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study robot foundation model scaling. The motivating bet is: Find embodied regimes where scaling data/model size predicts the wrong failure boundary.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For robot foundation model scaling, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2025a;b; 2026; 2025c; 2023; 2025d). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.560	0.440
observed_only_wam	hidden_mode_shift	0.185	0.815
augmentation_as_data_wam	nominal	0.640	0.360
augmentation_as_data_wam	hidden_mode_shift	0.263	0.738
scalar_uncertainty_wam	nominal	0.657	0.343
scalar_uncertainty_wam	hidden_mode_shift	0.307	0.693
proposed_branch_mechanism	nominal	0.802	0.198
proposed_branch_mechanism	hidden_mode_shift	0.578	0.422

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Scale-Law Exceptions in Robotics argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Scaling data-constrained language models, 2023. <https://doi.org/10.48550/arxiv.2305.16264>.
- Scaling cross-environment failure reasoning data for vision-language robotic manipulation, 2025a. <http://arxiv.org/abs/2512.01946v3>.
- Deepseq: High-throughput single-cell rna sequencing data labeling via web search-augmented agentic generative ai foundation models, 2025b. <http://arxiv.org/abs/2506.13817v1>.
- Transforming monolithic foundation models into embodied multi-agent architectures for human-robot collaboration, 2025c. <http://arxiv.org/abs/2512.00797v1>.
- Ecg-fm: an open electrocardiogram foundation model, 2025d. <https://doi.org/10.1093/jamiaopen/ooaf122>.
- Beyond data moore’s law: Towards sustainable scaling of foundation models, 2026. <https://doi.org/10.20944/preprints202511.1245.v1>.